

Simulated Volumetric Displays using Head Tracking

William Brown (whb2107)

April 14, 2011

My project seeks to create a program that simulates a volumetric display by tracking the user's head and modifying the image displayed on the screen to mimic parallax perspective.

1 Volumetric Displays

Volumetric displays (devices capable of three-dimensional video output) have long been relegated to the realm of science fiction; until recently, volumetric displays existed only in imaginations. In late 2009, Sony unveiled the world's first feasible volumetric display at the DC Expo in Tokyo [Zyga, 2009]. This device was capable of a resolution of 98 by 98 by 128-pixel resolution, and the cylindrical display was 13 centimeters in diameter. Today, two and a half years later, consumers are still waiting for volumetric displays to come to market, and while we are seeing a rise in 3D-television adoption, there seem to be very little demand for holographic or other kinds of volumetric displays.

On the other hand, simulated volumetric displays (or "parallax displays") like the one I plan on building have seen quite a few recent and not-so-recent implementations, with various success. Some rely on specialized hardware, such as Johnny Lee's implementation which uses a Wiimote as a camera and infrared LEDs mounted on glasses to determine the user's location in 3D space [Chung, 2008]. Other systems, mostly for using head motions as input to videogames, use tracking markers that are worn on a baseball cap. I was unable to find a popular implementation that did not utilize tracking markers, but there are published papers that suggest that it has been done, and with good results. I discuss this software and research in Section 3.

2 Motivation

I want to create a volumetric display because it has interesting, practical uses and is relatively simple to implement, given the libraries available. It also is something that is fun for an end-user to use, and easy to integrate into other software. I've found the available software either heavily platform-dependent (usually dependent on a Windows system), closed source, or difficult to use for anything other than

very basic, demonstrative purposes. If this project goes well, I would like to turn it into a cross-platform library for building parallax display into other 3D software after the end of the semester.

There are many possible use cases, but I thought of this project idea when working on my 3D printer. I spend lots of my time modelling objects for 3D printing, and having the ability to really “look” at an object on a convincing, three-dimensional display before it is printed would have saved me many reprints.

3 Relevant Software and Research

VR Displays using Wiimotes A similar project has been done by Johnny Lee, in which he displayed visual output that was modified based on head position to simulate a 3D display. However, he used a Wiimote to track the user’s head; the user wore glasses with infrared LEDs attached, and a Wiimote then communicated the location, rotation and distance of the glasses to the computer. He has released the source code for program he uses to do the head tracking [Chung, 2008]; it is written in C# and uses the Win32 API.

VR for Videogames: Cachya Cachya [Cachya Team, 2008] is a system designed to translate head location and pose into joystick or mouse events for use in videogames. It is primarily used as input for flight simulators, but could potentially be used in other games as well. Cachya uses a webcam system, but requires the user to wear a tracking marker on a hat. There is a very involved calibration step, as well, in which the user has to directly manipulate various image parameters (contrast, brightness, threshold, etc.) before they can use the system. Even after all of this calibration, a significant amount of smoothing is then done to prevent noisy input, which would be detrimental to a videogame. It uses machine learning techniques to attempt to identify the threshold between noise and input.

VR for Videogames: TrackIR Another major player in the videogame head tracking market is TrackIR [NaturalPoint Software, 2008]. This system uses a specialized camera and an marker device to track the user’s head. While the technology behind the camera is proprietary, it does seem to be an infrared camera, perhaps with some sort of structured light emission, as the marker device is reflective. It has six degrees of freedom, just like Cachya: roll, pitch, yaw, pan horizontally, pan vertically, and depth. Calibration is optional, and allows the user to change the response to various kinds of motion; for example, a user can select how sensitive the system should be to movements in different axes.

Head-Tracking Libraries The Enhanced Human-Computer Interface library [Baggio, 2010] is a library designed for real-time head tracking in three dimensions. It has six degrees of freedom, and outputs a 3-vector normal to the plane of the face. It was designed for a parallax display application, but then was changed

for more general-purpose usage; it now exists only as a system designed for adding head-based gestural control to other applications.

Review of Head-Tracking Methods An older paper, [Toyama, 1999], discusses various methods of face detection that his team tested in creating a parallax system not unlike my own. They tested five methods: foreground segmentation, color-based tracking, motion-based tracking, facial detection and a Bayesian network that simulates a combination of all tracking of the previously mentioned methods. They found that the motion-based system performed the best of any individual method, but the Bayesian network outperformed that method by a significant amount.

Review of Gesture Recognition Techniques A more recent paper, [C et al., 2010], proposes various algorithms for detecting hands, noses and eyes in images. The various methods of detecting eyes that are described could be very useful, as interocular distance could serve as an accurate measurement of depth. However, the outlined algorithms are also much more difficult to implement than just a standard face detection algorithm; if I find that facial detection is giving poor results, I will try some of the algorithms given here.

4 Implementation Details

I have two different methods in mind for how to implement face detection, both of which were proposed in [Toyama, 1999]. One is a segmentation based on color, and is not dissimilar from the segmentation method used in the first assignment. The other is a motion-based system, which evaluates the location of the face based on changes in pixel values. The color-based method assumes that the lighting of the face doesn't change and there are few skin-colored objects in the room, while the motion-based method assumes that nothing else in the frame is moving. I plan on implementing both of these methods to determine which gives me the best results, and then using that method going forward.

Once I have the bounding box of the face, I will then use the centroid of the face and it's size along the primary axis of the bounding rectangle to detect approximate relative location in three dimensions. After applying some smoothing, I will use the location in reference to the center of the frame (which I assume to be aligned with the center of the display) to determine the perspective from which the user would be viewing the scene.

At first, the system will be capable of some basic parallax using layers of images. This mode of operation will require the user to input various images and the "depth" at which each image lies. The location of the user's head will then be used to displace those images in relation to each other to create the appearance of multiple image planes with differing depth.

The next step will be to integrate the parallax system with OpenGL. Instead of a series of images, the system will read in a 3D scene from a descriptor file. The user moving their head will correspond to movements within this 3D scene.

5 Feynstein Integration

For my Programming Languages and Translators course, my team is designing a language for physical simulation called Feynstein. This language takes as input a description of a scene with objects and forces, and outputs a visual simulation of that scene in a window on the user's computer. If I have time after building my system as outlined above, I would like to integrate my volumetric display with the output of the Feynstein system. This would be completely optional, and this system integration would not be presented for a grade in Visual Interfaces.

6 System Evaluation

It is simple to come up with qualitative methods for judging the results of the system; after all, this is a system designed to simulate the physical phenomenon of parallax perspective. If the simulated third dimension is convincing to the user, and the image does appear to have depth, I would consider the project a success.

However, there are some qualitative measures I can use as well. The metric I will monitor is jitter in the perceived head location, which I anticipate will be a problem. I can log the location and depth of the face on each measurement, then fit a polynomial function to the results. If we assume that the user keeps their head relatively steady, the average error between the measured values and the polynomial will be equivalent to the error in the detected face location. This will be a good test of whether or not my facial detection and smoothing algorithms are working well.

7 Sample Output

Figures 1, 2, and 3 show some sketches of the sort of output I would expect. On the left of each image is a "webcam capture", representing where the user's face lies in relation to the scene. On the right is a sample rendering of how a cube would be displayed, if the user's face was in the corresponding location.

References

- D. Baggio. Enhanced human-computer interfaces through webcam image processing. <http://code.google.com/p/ehci/>, 2010.
- H. C. K. R. Shastry, M. Ravindran, M. V. V. N. S. Srikanth, and N. Lakshmikanth. Survey on Various Gesture Recognition Techniques for Interfacing Machines Based on Ambient Intelligence. *ArXiv e-prints*, November 2010.
- Cachya Team. Cachya head tracking software. <http://www.cachya.com/esight/overview.php>, 2008.
- J. L. Chung. Wii projects. <http://johnnylee.net/projects/wii/>, 2008.




Figure 1: The user's face is centered, so they see the front projection of the cube.




Figure 2: The user's face is on the left, so they see the cube from the right.




Figure 3: The face is in the top-right, so the cube is seen from the left and above.




Figure 4: When the user moves closer, the scene zooms in.

- NaturalPoint Software. TrackIR 5. <http://www.cachya.com/esight/overview.php>, 2008.
- K. Toyama. Head parallax tracking for control of a virtual space: a comparison of algorithms. In *Systems, Man, and Cybernetics, 1999. IEEE SMC '99 Conference Proceedings. 1999 IEEE International Conference on*, volume 6, pages 1 – 6 vol.6, 1999.
- L. Zyga. Sony unveils 360-degree 3d display. PhysOrg (<http://www.physorg.com/news175446089.html>), 2009.